(1) what dataset

- WildChat, a public dump of LLM prompts.

from datasets import load_dataset

Login using e.g. `huggingface-cli login` to access this dataset

ds = load_dataset("allenai/WildChat-4.8M")

(2) what types of insights

Prompt Topic Distribution — What are the most common themes users discuss with LLMs (e.g., coding, writing, reasoning, education)?

Conversation Length and Engagement — How long are typical conversations? Are certain topics associated with longer or more interactive dialogues?

Model Comparison — Are there stylistic or structural differences between responses from different models (e.g., GPT-4 vs. Claude vs. Gemini)?

Temporal Trends — How has user interaction evolved over time (e.g., before and after major model releases)?

(3) what the MCP server will expose

I will develop an **MCP server** that allows an LLM to "talk to" the WildChat dataset.
Planned endpoints:

- `search_conversations(keyword: str, limit: int)` — return sample conversations containing specific topics or phrases.
- `get_model_stats(model_name: str)` — summarize dataset statistics for a given LLM (average length, sentiment, etc.).
- `compare_models(model_a: str, model_b: str, metric: str)` — compare two models on length, politeness, or response diversity.
- `get_top_topics(n: int)` — return top N most frequent conversation categories.

These endpoints will let a connected LLM (e.g., Claude or ChatGPT) perform follow-up reasoning, such as:

> "Show me how user tone differs when talking to GPT-4 versus Gemini,"
> or
> "Retrieve examples of educational prompts involving mathematics."